

Advanced Machine Learning, Assignment 2

Youssef Taoudi, yousseft@kth.se

January 7, 2020

2.1 Knowing The Rules

Question 1

Yes.

Question 2

No collaborations (aside from discussions in the slack group)

Question 3

No.

2.2 Dependencies in a DGM

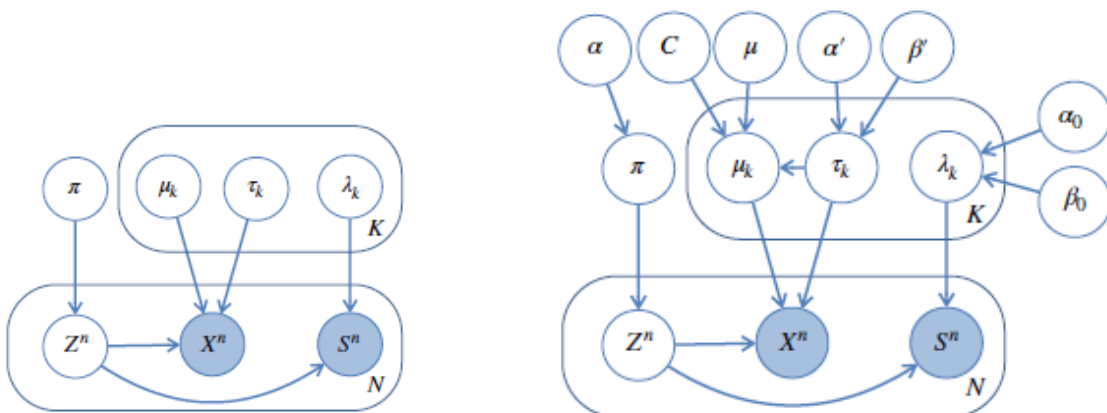


Figure 1: Graphical Models of Figure 1 and Figure 2 from the assignment

Question 4

Yes

Question 5

No

Question 6

Yes

Question 7

No

Question 8

No

Question 9

No

2.3 Tree GM

Question 10

The probability $p(\beta|T, \Theta)$ is the product of the probability of the observed value of a leaf given the observed tree as starting on any particular leaf will ensure traversal through the entire tree.

$$p(\beta|T, \Theta) = p(X_L = V_L|X_o) \quad (1)$$

Where X_L is a leaf node, V_L is the corresponding value to the node and X_o are all observations in the tree.

$$p(X_L = V_L|X_o) = p(X_L = V_L|X_{o \uparrow L})p(X_L = V_L|X_{o \downarrow L}) \quad (2)$$

In this case $X_{o \downarrow L}$ are all nodes below X_L and $X_{o \uparrow L}$ is the rest. The expression can be rewritten to:

$$p(X_L = V_L|X_o) = s(X_L, V_L)t(X_L, V_L) \quad (3)$$

Starting with $s(X_L, V_L)$, it is a top-down recursive function that calculates $s(X_L, V_L)$ sums over all possible categories for the child nodes [4](#). The base case is at the leaves where

$s(X_L, V_L)$ returns 1 if the observed leaf value V_L equals input value i 5. Note that the CPD is also calculated for the children over all categories.

$$s(X_u, i) = \sum_j^K p(X_{child1} = j | X_u = i) s(X_{child2}, j) \sum_j^K p(X_{child1} = j | X_u = i) s(X_{child2}, j) \quad (4)$$

$$s(X_L, i) = \begin{cases} 1, & \text{if } V_L = i \\ 0, & \text{otherwise} \end{cases} \quad (5)$$

$t(X_u, i)$ on the other hand is a bottom-up approach that visits all parent and children nodes 6. Once again, function t will sum over all combinations of parent/sibling category pairs. In each step the CPD of the current node X_u and sibling node is calculated.

$$t(X_u, i) = \sum_{j,k}^K t(X_{parent}, j) p(X_u = i | X_{parent} = j) p(X_{sibling} = k | X_{parent} = j) s(X_{sibling}, k) \quad (6)$$

Once a given $t(X_u, i)$ or $s(X_u, i)$ has been calculated, it is stored in a directory and can be accessed in constant time if needed without passing through the tree again. Code for the algorithm is included in the appendix A.

Question 11

	Small Tree	Medium Tree	Large Tree
Sample 1	0.008753221441670067	8.66416414170832e-17	1.2296785012112113e-65
Sample 2	0.03839692509792911	5.394284454090607e-18	1.4347770777980813e-63
Sample 3	0.009129106859990063	8.892415333536362e-18	3.0954910161498576e-66
Sample 4	0.0214406975419561	1.1222302136292958e-18	3.4231977224272667e-69
Sample 5	0.011945567814215125	7.589341572491351e-19	4.822393947666209e-67

2.4 Simple VI

Question 12

Following the derivations from Bishop [1, pp. 470–471] and Murphy[2, pp. 742–744] by inferring two factors $q(\tau)$ and $q(\mu)$ two derive the posterior. The factors are derived by inferring four hyperparameters α_N, β_N, μ_N and λ_N iteratively from equations 7 - 8. Note that since β_N is needed to calculate λ_N and vice versa, the first β_N that is used is a guessed value. Note that $\alpha_0 = \beta_0 = \lambda_0 = \mu_0 = 0$ (Initial guess). An example of how the guessed posterior converges over several iterations can be seen in figure 2 and code can be found in Appendix B.

$$\mu_N = \frac{\lambda_0 \mu_0 + N \bar{x}}{\lambda_0 + N}, \lambda_N = (\lambda_0 + N) \frac{\alpha_N}{\beta_N} \quad (7)$$

$$\alpha_N = \alpha_0 + \frac{N+1}{2},$$

$$\beta_N = \beta_0 + \lambda_0 (\lambda_N + \mu_N^2 + \mu_0^2 - 2\mu_N\mu_0) + \frac{1}{2} \sum_{i=1}^N (x_i^2 + \lambda_N + \mu_N^2 - 2\mu_N x_i) \quad (8)$$

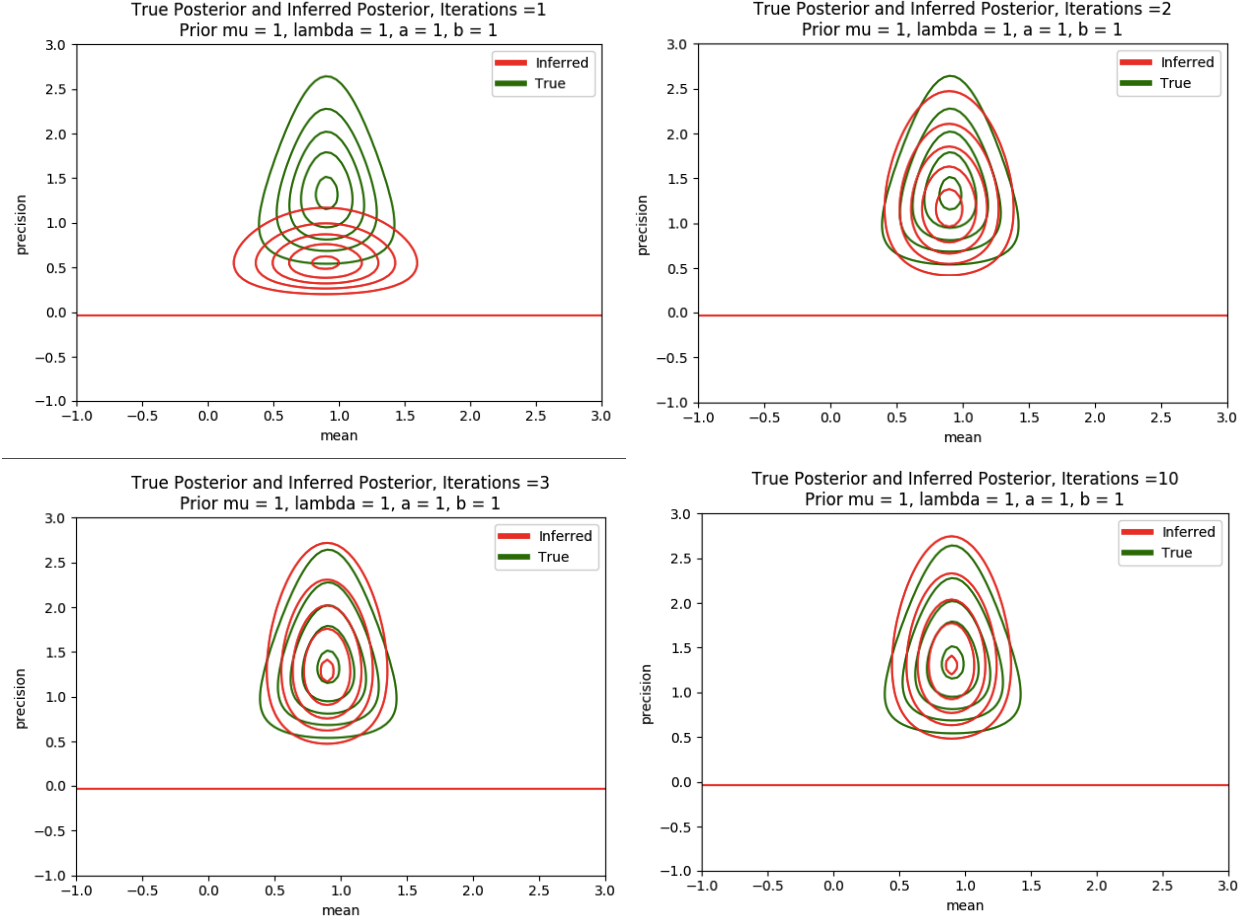


Figure 2: Plot showcasing how the inferred posterior converges after multiple iterations.

Question 13

The exact posterior, $P(\mu, \tau|D) = P(\mu|\tau)P(\tau)P(D|\mu, \tau)$. $P(\mu|\tau)$ is Gaussian while $P(\tau)$ is Gamma distributed, meaning the conjugate priors will result in a Normal-Gamma distribution which is normalized by the likelihood $P(\mu, \tau|D)$ to form the exact posterior. [2, p. 742]. Examples of exact posteriors can be showcased in figures 2 and 3 where the exact posterior are drawn in green. Note, $p(\mu|\tau) \sim \mathcal{N}(\mu_0, (\lambda\tau)^{-1})$ and $p(\tau) \sim \text{Gamma}(\alpha, \beta)$.

Question 14

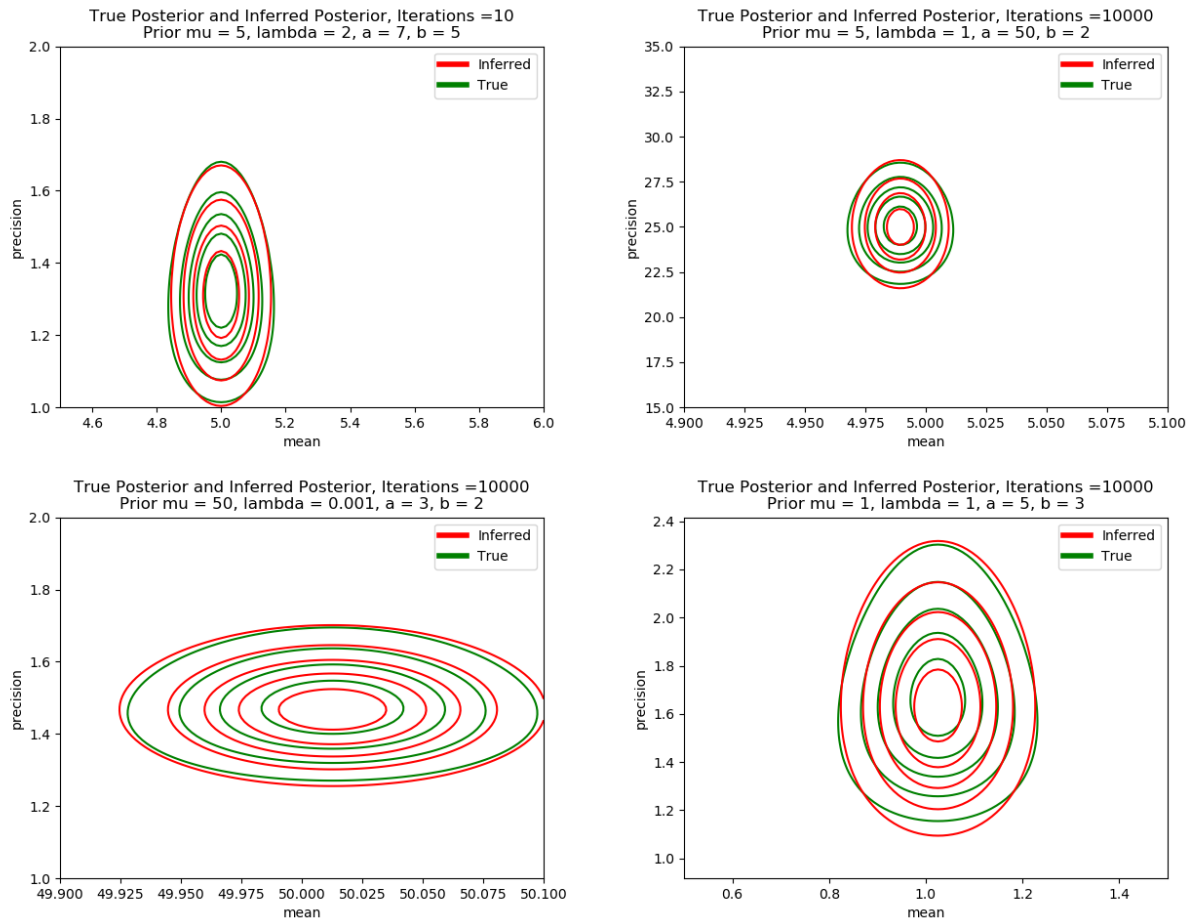


Figure 3: Plots of the true posterior compared to the inferred posterior for 4 different cases with different values on the hyperparameters.

As seen in figure 3 the converged posterior takes different shapes depending on hyperparameters but it seems to match the exact posterior well most of the time.

Appendices

A: TreeGM Algorithm

```
# Calculating s values for the tree and storing them
def s_root(tree_topology, theta, beta):
    prob = 0
    s_dict = defaultdict(dict)

    def S(u, j, children):
        if s_dict[u].get(j) is not None:
            return s_dict[u][j]
        if len(children) < 1:
            if beta.astype(int)[u] == j:
                s_dict[u][j] = 1
                return 1
            else:
                s_dict[u][j] = 0
                return 0
        result = np.zeros(len(children))
        for child_nr, child in enumerate(children):
            for category in range(0, len(theta[0])):
                result[child_nr] += S(child, category, find_children(child,
                    tree_topology, beta)) * CPD(theta, child,

        s_result = np.prod(result)
        s_dict[u][j] = s_result
        return s_result

    for i, th in enumerate(theta[0]):
        prob += S(0, i, find_children(0, tree_topology, beta)) * CPD(theta, 0, i)
    return s_dict

def calculate_likelihood(tree_topology, theta, beta):
    """
    This function calculates the likelihood of a sample of leaves.
    :param: tree_topology: A tree topology. Type: numpy array. Dimensions:
        (num_nodes, )
    :param: theta: CPD of the tree. Type: numpy array. Dimensions: (num_nodes, K)
    :param: beta: A list of node assignments. Type: numpy array. Dimensions:
        (num_nodes, )
    Note: Inner nodes are assigned to np.nan. The leaves have values in [K]
    :return: likelihood: The likelihood of beta. Type: float.
    """
```

```

s_dict = s_root(tree_topology, theta, beta)
t_dict = defaultdict(dict)

"""Recursively calculates t(u,i) if it has not already been calculated"""
def t(u, i, parent, sibling):
    if t_dict[u].get(i) is not None: # If it has already been calculated
        return t_dict[u][i]

    if np.isnan(parent): # If root
        return CPD(theta, u, i)
    if sibling is None: # If no siblings
        result = 0
        for j in range(0, len(theta[0])):
            result += CPD(theta, u, i, j) * t(parent, j, tree_topology[parent],
                                                find_sibling(parent, tree_topology))

        t_dict[u][i] = result
        return result

    parent = int(parent)
    result = 0
    for j in range(0, len(theta[0])):
        for k in range(0, len(theta[0])):
            result += CPD(theta, u, i, j) * CPD(theta, sibling, k, j) * \
                s_dict[sibling][k] * t(parent, j, tree_topology[parent],
                                        find_sibling(parent, tree_topology))
    t_dict[u][i] = result
    return result

""" Find the first leaf and calculate its likelihood """
for leaf, cat in enumerate(beta):
    if not np.isnan(cat):
        return t(leaf, cat, int(tree_topology[leaf]),
                find_sibling(leaf, tree_topology)) * s_dict[leaf][cat]

```

B: Variational Inference Algorithm

```

def expected_mu(lamb0, X, mu0, mu_n, lamb_n):
    E_mu2 = lamb_n ** (-1) + mu_n ** 2
    square_sum = np.sum((X ** 2) - (2 * X * mu_n) + E_mu2)
    return (1 / 2 * square_sum) + lamb0 * ((mu0 ** 2) - (2 * mu0 * mu_n) + E_mu2)

def approx_a(a0, n):
    return a0 + ((n + 1) / 2)

#Note, argument l0 = Lambda_0, NOT TEN
def approx_mu(l0, m0, X, n):

```

```

    return (l0 * m0 + n * np.average(X)) / (l0 + n)

def approx_lambda(l0, a_n, b_n, n):
    return (l0 + n) * (a_n / b_n)

def approx_b(m0, m_n, l_n, l0, b0):
    return b0 + expected_mu(l0, X, m0, m_n, l_n)

def VariationalInference(mu0, lamb0, a0, b0, X, iterations):
    i = 0
    la = 1 #Initial guess
    be = 1 #Initial guess
    mu = mu0
    al = a0
    while i < iterations:
        al = approx_a(a0, N)
        mu = approx_mu(lamb0, mu0, X, N)
        be = approx_b(mu0, mu, la, lamb0, b0)
        la = approx_lambda(lamb0, al, be, N)
        i += 1
    if i == iterations:
        return mu, la, al, be

```

References

- [1] Christopher M. Bishop. *Pattern recognition and machine learning*. en. Information science and statistics. New York: Springer, 2006. ISBN: 978-0-387-31073-2.
- [2] Kevin P. Murphy. *Machine learning: a probabilistic perspective*. en. Adaptive computation and machine learning series. Cambridge, MA: MIT Press, 2012. ISBN: 978-0-262-01802-9.